

<https://github.com/UCR-HPC/cs211-hw3-sieving-prime-HugoWan0504>

Q1

Screenshot of the sieve0's test outputs:

```
func=sieve0, pnum=32, count=455052511, time=17.709125
func=sieve0, pnum=64, count=455052511, time=8.972398
func=sieve0, pnum=96, count=455052511, time=5.978825
func=sieve0, pnum=128, count=455052511, time=4.495173
func=sieve0, pnum=160, count=455052511, time=3.570440
```

Organized Table of **function sieve0** and **prime count = 455052511**:

Cores	Time (seconds)
32	17.709125
64	8.972398
96	5.978825
128	4.495173
160	3.570440

Q2

Screenshot of the sieve1's test outputs:

```
func=sieve1, pnum=32, count=455052511, time=5.586815
func=sieve1, pnum=64, count=455052511, time=3.348253
func=sieve1, pnum=96, count=455052511, time=2.320773
func=sieve1, pnum=128, count=455052511, time=1.721402
func=sieve1, pnum=160, count=455052511, time=1.385505
```

Organized Table of **function sieve1** and **prime count = 455052511**:

Cores	Time (seconds)
32	5.586815
64	3.348253
96	2.320773

128	1.721402
160	1.385505

Sieve1.c codes:

```

1  #ifndef __SIEVE1_C__
2  #define __SIEVE1_C__
3
4  #include "include.h"
5
6  // Follow instructions based on HW3 guide. Thanks, TA.
7  void sieve1(unsigned long long *global_count, unsigned long long n, int pnum, int pid)
8  {
9      unsigned long long low_value = 3 + 2 * ((pid * ((n - 3) / 2 + 1)) / pnum); // the smallest value handled by this process
10     unsigned long long high_value = 3 + 2 * (((pid + 1) * ((n - 3) / 2 + 1)) / pnum) - 2; // the largest value handled by this process
11     unsigned long long size = (high_value - low_value) / 2 + 1; // number of integers handled by this process
12
13     if (1 + (n - 1) / pnum < (int)sqrt((double)n)) // high_value of process 0 should be larger than floor(sqrt(n))
14     {
15         if (pid == 0)
16             printf("Error: Too many processes.\n");
17         MPI_Finalize();
18         exit(0);
19     }
20
21     char *marked = (char*)malloc(size); // array for marking multiples. 1 means multiple and 0 means prime
22     if (marked == NULL)
23     {
24         printf("Error: Cannot allocate enough memory.\n");
25         MPI_Finalize();
26         exit(0);
27     }
28     memset(marked, 0, size);
29
30     unsigned long long index = 0; // index of current prime among all primes (only works for process 0)
31     unsigned long long prime = 3; // current prime broadcasted by process 0
32     do {
33         unsigned long long first; // index of the first multiple among values handled by this process
34         if (prime * prime > low_value)
35             first = ((prime * prime) - low_value) / 2;
36         else if (low_value % prime == 0)
37             first = 0;
38         else
39             first = ((prime * (((low_value / prime) + 1) / 2) * 2 + 1)) - low_value) / 2;
40
41         for (unsigned long long i = first; i < size; i += prime)
42             marked[i] = 1;
43
44         index++;
45         if (pid == 0)
46         {
47             while (marked[index] == 1)
48                 index++;
49             prime = (index * 2) + 3;
50         }
51         MPI_Bcast(&prime, 1, MPI_INT, 0, MPI_COMM_WORLD);
52     } while (prime * prime <= n);
53
54     unsigned long long count = 0; // local count of primes
55     for (unsigned long long i = 0; i < size; i++)
56         if (marked[i] == 0)
57             count++;
58
59     if (pid == 0) count++; // add prime 2 to the count
60     MPI_Reduce(&count, global_count, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);
61
62     // free memory
63     free(marked);
64 }
65
66
67 #endif

```

Q3

Screenshot of the sieve2's test outputs:

```
func=sieve2, pnum=32, count=455052511, time=5.639904
func=sieve2, pnum=64, count=455052511, time=3.297908
func=sieve2, pnum=96, count=455052511, time=2.205112
func=sieve2, pnum=128, count=455052511, time=1.745691
func=sieve2, pnum=160, count=455052511, time=1.469822
```

Organized Table of **function sieve2** and **prime count = 455052511**:

Cores	Time (seconds)
32	5.639904
64	3.297908
96	2.205112
128	1.745691
160	1.469822

Sieve2.c codes:

```
1  #ifndef __SIEVE2_C__
2  #define __SIEVE2_C__
3
4  #include "include.h"
5
6  // Follow instructions based on HW3 guide. Thanks, TA.
7  void sieve2(unsigned long long *global_count, unsigned long n, int pnum, int pid)
8  {
9      unsigned long long low_value = 3 + 2 * ((pid * ((n - 3) / 2 + 1)) / pnum); // smallest value handled by this process
10     unsigned long long high_value = 3 + 2 * (((pid + 1) * ((n - 3) / 2 + 1)) / pnum) - 2; // largest value handled by this process
11     unsigned long long size = (high_value - low_value) / 2 + 1; // number of integers handled by this process
12
13     if ((1 + (n - 1)) / pnum < (int)sqrt((double)n)) // high_value of process 0 should be larger than floor(sqrt(n))
14     {
15         if (pid == 0)
16             printf("Error: Too many processes.\n");
17         MPI_Finalize();
18         exit(0);
19     }
20
21     char *marked = (char*)malloc(size); // array for marking multiples. 1 means multiple and 0 means prime
22     if (marked == NULL)
23     {
24         printf("Error: Cannot allocate enough memory.\n");
25         MPI_Finalize();
26         exit(0);
27     }
28     memset(marked, 0, size);
```

```

30     unsigned long long sqrt_n = (unsigned long long)sqrt((double)n);
31     unsigned long long sieve_size = (sqrt_n - 3) / 2 + 1; // size for local sieving primes
32     char *sieve = (char*)malloc(sieve_size); // array for local sieving primes
33     if (sieve == NULL)
34     {
35         printf("Error: Cannot allocate enough memory for sieve.\n");
36         MPI_Finalize();
37         exit(0);
38     }
39     memset(sieve, 0, sieve_size);
40
41     // Generate sieving primes locally
42     for (unsigned long long i = 0; i < sieve_size; i++)
43     {
44         if (sieve[i] == 0) // Found a prime
45         {
46             unsigned long long prime = 2 * i + 3;
47             for (unsigned long long j = (prime * prime - 3) / 2; j < sieve_size; j += prime)
48                 sieve[j] = 1;
49         }
50     }
51
52     // Use sieving primes to mark composites in the range
53     for (unsigned long long i = 0; i < sieve_size; i++)
54     {
55         if (sieve[i] == 0) // Found a sieving prime
56         [
57             unsigned long long prime = 2 * i + 3;
58             unsigned long long first; // First multiple to mark
59             if (prime * prime > low_value)
60                 first = ((prime * prime) - low_value) / 2;
61             else if (low_value % prime == 0)
62                 first = 0;
63             else
64                 first = (((low_value / prime) + 1) / 2) * 2 + 1;
65
66             for (unsigned long long j = first; j < size; j += prime)
67                 marked[j] = 1;
68         ]
69     }
70
71     unsigned long long count = 0; // Local count of primes
72     for (unsigned long long i = 0; i < size; i++)
73         if (marked[i] == 0)
74             count++;
75
76     if (pid == 0) count++; // Add prime 2 to the count
77     MPI_Reduce(&count, global_count, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);
78
79     // Free allocated memory
80     free(marked);
81     free(sieve);
82 }
83
84 #endif

```

Q4

Screenshot of the sieve3's test outputs:

```
|func=sieve3, pnum=32, count=455052511, time=1.279721
|func=sieve3, pnum=64, count=455052511, time=0.635969
|func=sieve3, pnum=96, count=455052511, time=0.424961
|func=sieve3, pnum=128, count=455052511, time=0.318910
|func=sieve3, pnum=160, count=455052511, time=0.253828
```

Organized Table of **function sieve3** and **prime count = 455052511**:

Cores	Time (seconds)
32	1.279721
64	0.635969
96	0.424961
128	0.318910
160	0.253828

Sieve3.c codes:

```
1 #ifndef __SIEVE3_C__
2 #define __SIEVE3_C__
3
4 #include "include.h"
5
6 // Follow instructions based on HW3 guide. Thanks, TA.
7 #define BLOCK_SIZE 1024 * 32 // Based on cache optimization hints
8
9 void sieve3(unsigned long long *global_count, unsigned long long n, int pnum, int pid)
10 {
11     unsigned long long sqrt_n = (unsigned long long)sqrt((double)n);
12
13     // Step 1: calculate local primes up to sqrt(n)
14     unsigned long long local_sieve_size = (sqrt_n - 3) / 2 + 1; // Only odd numbers up to sqrt(n)
15     char *local_sieve = (char *)malloc(local_sieve_size);
16     if (!local_sieve)
17     {
18         printf("Error: Unable to allocate memory for local sieve array.\n");
19         MPI_Finalize();
20         exit(0);
21     }
22     memset(local_sieve, 0, local_sieve_size);
23
24     // Generate local primes
25     for (unsigned long long i = 0; i < local_sieve_size; i++)
26     {
27         if (local_sieve[i] == 0)
28         {
29             unsigned long long prime = 2 * i + 3;
30             for (unsigned long long j = (prime * prime - 3) / 2; j < local_sieve_size; j += prime)
31                 local_sieve[j] = 1;
32         }
33     }
34 }
```

```

35 // Store local primes for marking later
36 unsigned long long *local_primes = (unsigned long long *)malloc(local_sieve_size * sizeof(unsigned long long));
37 unsigned long long local_prime_count = 0;
38 for (unsigned long long i = 0; i < local_sieve_size; i++)
39 {
40     if (local_sieve[i] == 0)
41         local_primes[local_prime_count++] = 2 * i + 3;
42 }
43 free(local_sieve);
44
45 // Step 2: Define the range for this process
46 unsigned long long low_value = 3 + 2 * ((pid * ((n - 3) / 2 + 1)) / pnum);
47 unsigned long long high_value = 3 + 2 * (((pid + 1) * ((n - 3) / 2 + 1)) / pnum) - 2;
48 unsigned long long size = (high_value - low_value) / 2 + 1;
49
50 // Allocate memory for marking composites
51 char *marked = (char *)malloc(BLOCK_SIZE);
52 if (!marked)
53 {
54     printf("Error: Unable to allocate memory for marked array.\n");
55     MPI_Finalize();
56     exit(0);
57 }
58
59 unsigned long long count = 0;
60
61 // Step 3: Block processing for cache efficiency
62 for (unsigned long long block_low = low_value; block_low <= high_value; block_low += 2 * BLOCK_SIZE)
63 {
64     unsigned long long high_value = block_low + BLOCK_SIZE - 1;
65     if (block_high > high_value)
66         block_high = high_value;
67
68     unsigned long long block_size = (block_high - block_low) / 2 + 1;
69     memset(marked, 0, block_size);
70
71     // Mark multiples of local primes in the current block
72     for (unsigned long long i = 0; i < local_prime_count; i++)
73     {
74         unsigned long long prime = local_primes[i];
75         unsigned long long first;
76
77         if (prime * prime > block_low)
78             first = (prime * prime - block_low) / 2;
79         else if (block_low % prime == 0)
80             first = 0;
81         else
82         {
83             unsigned long long temp = (block_low + prime) / prime;
84             if (temp % 2 == 0)
85                 temp++;
86             first = (temp * prime - block_low) / 2;
87         }
88
89         for (unsigned long long j = first; j < block_size; j += prime)
90             marked[j] = 1;
91     }
92
93     // Count primes in the current block
94     for (unsigned long long i = 0; i < block_size; i++)
95         if (marked[i] == 0)
96             count++;
97 }
98
99 // Include prime 2 in the count
100 if (pid == 0) count++;
101
102 // Step 4: Reduce the counts across all processes
103 MPI_Reduce(&count, global_count, 1, MPI_UNSIGNED_LONG_LONG, MPI_SUM, 0, MPI_COMM_WORLD);
104
105 free(marked);
106 free(local_primes);
107 }
108
109 #endif

```

Summary Tables of Execution Time (in seconds) & Speedup:

Execution Time Comparison:

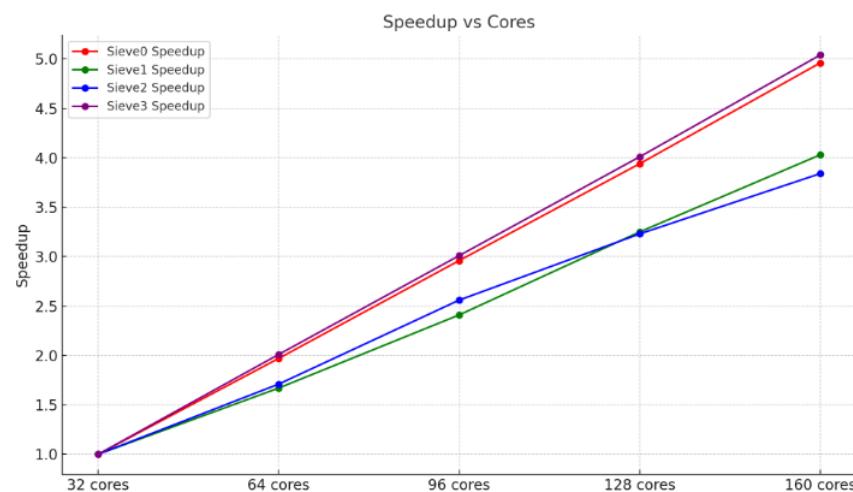
Cores	Sieve0 Time	Sieve1 Time	Sieve2 Time	Sieve3 Time
32	17.709125	5.586815	5.639904	1.279721
64	8.972398	3.348253	3.297908	0.635969
96	5.978825	2.320773	2.205112	0.424961
128	4.495173	1.721402	1.745691	0.318910
160	3.570440	1.385505	1.469822	0.253828

Speedup Comparison:

Cores	Sieve0 Speedup	Sieve1 Speedup	Sieve2 Speedup	Sieve3 Speedup
32	1.00	1.00	1.00	1.00
64	1.97	1.67	1.71	2.01
96	2.96	2.41	2.56	3.01
128	3.94	3.25	3.23	4.01
160	4.96	4.03	3.84	5.04

Slope of Speedups Comparison:

Sieve Version	Slope of Speedups
Sieve0	0.0309
Sieve1	0.0239
Sieve2	0.0225
Sieve3	0.0315



In conclusion:

Sieve0 is the best when it comes to scaling with more cores, making it a great choice for larger parallel computing tasks. Sieve3 works really well with fewer cores but doesn't improve as much when more cores are added, so it might need some tweaks to handle bigger systems. Sieve1 and Sieve2 give decent performance overall, but they are not as scalable as Sieve0 or as fast as Sieve3 on smaller setups. In short, Sieve0 is great for scalability, Sieve3 is better for smaller systems, and Sieve1 and Sieve2 are balanced but could be improved.